Enyi (Olivia) Jiang

AI Safety & Alignment

enyij@stanford.edu | 217-200-2424

enyijiang.github.io | github.com/enyijiang | LinkedIn

Profile

AI safety and alignment researcher and final-year UIUC CS Ph.D. candidate (expected 2027), visiting Stanford STAIR. Develop representation-level safety evaluation, preference optimization, and selective intervention methods that preserve model utility. First/co-first-authored work at NeurIPS, ICLR, and TMLR, with foundations in adversarial robustness and domain adaptation. At Meta, built image/video generation methods for production pipelines and ran multi-node, multi-GPU training with FSDP.

Technical Expertise

Research Methods	Representation-level safety evaluation, adversarial robustness, preference optimization, selective safety adapters, diffusion distillation, domain adaptation.
Training & Evaluation	PyTorch, TRL, FSDP, multi-node/multi-GPU training, DPO/SimPO, adversarial testing, controlled experiments, ablation studies.
Exploratory Work	GRPO for tool-use tasks; agent monitoring/steering, with implemented code and training experiments; agent rollouts and multi-turn trajectories.
Programming & Tools	Python, NumPy, pandas, scikit-learn, C/C++, Java, Linux, Bash, Git.

Education

University of Illinois Urbana-Champaign Ph.D. expected 2027

Ph.D., Computer Science (M.S. 2023).

UIUC / Zhejiang University, ZJU-UIUC Institute 2016 – 2020

B.S., Computer Engineering (UIUC, Highest Honors); B.Eng., Electronics and Computer Engineering (ZJU).

Experience

Stanford University 2026 – Present | Stanford, CA

Visiting Student Researcher, Stanford Trustworthy AI Research (STAIR) Lab

- Co-led **CLEAR**, using a hidden-state gate to control a safety LoRA adapter. Reduced HarmBench attack success from **32.3% to 0.5%** on Llama-3-8B-Instruct, with up to **7.1 percentage points higher GSM8K accuracy** than globally applied safety SFT or LoRA.
- Developed representation-level safety auditing and a Latent Vulnerability Score. Across three model families, constructed models showed **2.5–3.1× the vulnerability** of aligned bases at the targeted middle layer despite safe-looking outputs.

Meta Summers 2025 & 2026 | Menlo Park, CA

Research Intern | Image & Video Generation

- **2025 | Diffusion acceleration:** Implemented and evaluated **DMD1/DMD2** for video-generation diffusion models, reducing denoising from **50 steps to 4** with a quality tradeoff.
- **2025 | Training engineering:** Ran **multi-node, multi-GPU PyTorch FSDP training**; debugged data quality, hyperparameters, GPU memory, and gradient flow across alternating model updates.
- **2025 | Production impact:** Contributed the distillation implementation and experimental validation that colleagues subsequently extended into Meta's **production pipeline**. Received an **internship return offer**.
- **2026 | Rapid learning & implementation:** Rapidly learned and implemented **drifting models and AnyFlow** for image/video generation, evaluating **SFT and distillation recipes** and obtaining preliminary AnyFlow results.

Argonne National Laboratory Jun – Aug 2023 | Remote

Research Intern

- Integrated federated domain adaptation algorithms into open-source APPFL; reproduced results and implemented custom losses and evaluation metrics with the team.

TensorOpera AI (formerly FedML) Feb – May 2023 | Remote

Research Intern

- Developed federated learning algorithms across four hospital silos; explored diffusion models and self-supervised learning with limited data.

Tencent AI Lab Dec 2020 – Apr 2021 | Guangdong, China

NLP Intern

- Designed Chinese word segmentation labeling standards and a benchmark spanning difficult cases and diverse domains; cleaned word embeddings and first-authored the associated patent application.

Research: AI Safety, Alignment & RL

- **Preference optimization | GANPO:** Designed and implemented a latent adversarial regularizer aligning policy/reference representations; integrated with DPO/SimPO using **TRL and UltraFeedback**. Compared against baselines with matched epochs and similar hyperparameters. On Llama-3-8B-Instruct, improved SimPO's AlpacaEval 2.0 length-controlled win rate from **48.31% to 50.48%**, with comparable mean response length (1,836 vs. 1,834). [Code]
- **Reasoning safety | MATCHA:** Built answer-conditioned token- and embedding-level adversarial probes and an automated evaluation framework. Showed correct answers can coexist with inconsistent reasoning; identified fragility in multi-step and commonsense tasks and studied transfer to black-box LLMs. [Code]
- **Robustness foundations | RAMP:** Framed multi-norm adversarial training as domain adaptation and transferred gradient projection from earlier federated-learning research, combining clean and adversarial gradients to improve the empirical clean/robustness tradeoff. Achieved **44.6% union robust accuracy with 81.2% clean accuracy** on CIFAR-10 (PreActResNet-18, trained from scratch, AutoAttack). First author, NeurIPS 2024.
- **Certified robustness | CURE:** Extended multi-norm robustness to certified training through bound alignment and training strategies. First author, TMLR 2026.
- **Agent training | Ongoing work:** Implemented code and ran training experiments for GRPO-based tool use and agent monitoring/steering.

Research: Applied AI

- **Climate AI | PyroAdapt:** Built a pretrain–retrieve–rank framework for wildfire prediction under spatial heterogeneity and temporal shift. Across California, raised daily average precision from **21.62% to 24.57%** and Top-5% recall from **18.70% to 22.79%**, capturing **344 additional positive cell-days** under a fixed detection budget.
- **Healthcare AI | Federated domain adaptation:** Designed gradient projection and auto-weighted aggregation to combine source and target model updates under domain shift. Improved adaptation with limited target data across healthcare and general datasets; studied demographic information as auxiliary signals for COVID-19 imaging in MIDRC.

Selected Publications

* denotes equal contribution.

- **Enyi Jiang**, Changming Xu, Nischay Singh, Tian Qiu, Gagandeep Singh. *Robust Answers, Fragile Logic: Probing the Decoupling Hypothesis in LLM Reasoning.* (TMLR 2026)
- **Enyi Jiang**, David S. Cheung, Gagandeep Singh. *Towards Generalized Certified Robustness with Multi-Norm Training.* (TMLR 2026)
- **Enyi Jiang**, Gagandeep Singh. *RAMP: Boosting Adversarial Robustness Against Multiple ℓ_p Perturbations for Universal Robustness.* (NeurIPS 2024)
- **Enyi Jiang***, Yibo Jacky Zhang*, Oluwasanmi O. Koyejo. *Principled Federated Domain Adaptation: Gradient Projection and Auto-Weighting.* (ICLR 2024)
- Chengxiao Wang*, **Enyi Jiang***, Xiaojing Liao, Sanmi Koyejo. *CLEAR: Continuous Latent Adapter Routing for Utility-Preserving LLM Safety Alignment.* Trustworthy AI for Good Workshop @ NeurIPS 2026 (accepted). (Paper)
- **Enyi Jiang***, Anders Gjølbye*, Yibo Jacky Zhang, Sanmi Koyejo. *When Behavioral Safety Evaluation Fails: A Representation-Level Perspective.* Preprint, 2026. Stanford Causal Science Conference, Frontiers in AI Evaluation (poster). (Paper) (Poster)
- **Enyi Jiang**, Yibo Jacky Zhang, Yinglun Xu, Andreas Haupt, Nancy Amato, Sanmi Koyejo. *Latent Adversarial Regularization for Offline Preference Optimization.* Preprint, 2026. (Paper)

Full publication list on Google Scholar.

Service

Founder, Inside AI Safety, Inc. — independent nonprofit AI safety research organization.
Reviewer: NeurIPS, ICLR, ICML, TMLR, CVPR, ECCV, AAAI, AISTATS, COLM, TPAMI.
Mentor, AI Alignment @ UIUC: Red-teaming/Adversarial and Interpretability streams.