Enyi (Olivia) Jiang

AI Safety & Alignment | RL for LLMs & Agents | Multimodal Generation

enyij@stanford.edu | enyij2@illinois.edu | 217-200-2424

enyijiang.github.io | github.com/enyijiang | LinkedIn

Profile

AI safety and alignment researcher with experience in LLM safety evaluation, preference optimization, adversarial robustness, and multimodal generation. Proficient in GRPO and RLHF, with hands-on experience in agent rollouts, tool-use, and multi-turn trajectory training. Develop methods to expose hidden model vulnerabilities and improve safety while preserving useful capabilities. Research experience at Stanford STAIR and Meta; final-year CS Ph.D. candidate at UIUC (expected 2027).

Technical Expertise

RL & Agent Training	GRPO, RLHF, agent rollouts, tool-use, multi-turn trajectory training.
Safety & Alignment	Preference optimization (DPO, SimPO), latent adversarial regularization, representation-level safety evaluation, adversarial testing, selective safety adapters.
Multimodal Models	Image/video generation, diffusion models, distillation, generation acceleration.
Engineering	Python, PyTorch, NumPy, pandas, scikit-learn, C/C++, Java, Linux, Bash, Git.

Experience

Stanford University 2026 – Present | Stanford, CA

Visiting Researcher, Stanford Trustworthy AI Research (STAIR) Lab

- Develop LLM safety evaluation and alignment methods using internal representations to detect vulnerabilities, improve preference learning, and selectively intervene on unsafe behavior.
- Lead or co-lead representation-level safety auditing, latent adversarial regularization, and utility-preserving safety adaptation projects.

Meta Summers 2025 & 2026 | Menlo Park, CA

Research Intern | Multimodal Image & Video Generation

- Developed diffusion distillation and acceleration solutions for Meta's image and video generation models, focusing on generation quality and computational efficiency.
- Investigated and compared diffusion distillation approaches to identify quality–efficiency tradeoffs for multimodal generation.

Argonne National Laboratory Jun – Aug 2023 | Remote

Research Intern

- Integrated federated domain adaptation algorithms into open-source APPFL; reproduced results and implemented custom losses and evaluation metrics with the team.

TensorOpera AI (formerly FedML) Feb – May 2023 | Remote

Research Intern

- Developed federated learning algorithms for healthcare across four hospital silos; explored diffusion models and self-supervised learning for limited-data settings.

Tencent AI Lab Dec 2020 – Apr 2021 | Guangdong, China

NLP Intern

- Designed Chinese word segmentation labeling standards and a benchmark spanning difficult cases and diverse domains; cleaned word embeddings and first-authored the associated patent application.

Education

University of Illinois Urbana-Champaign 2023 – 2027 (expected)

Ph.D., Computer Science.

University of Illinois Urbana-Champaign 2020 – 2023

M.S., Computer Science.

UIUC / Zhejiang University, ZJU-UIUC Institute 2016 – 2020

B.S., Computer Engineering (UIUC, Highest Honors); B.Eng., Electronics and Computer Engineering (ZJU).

AI Safety & Alignment: Selected Work

- **Safety auditing:** Developed an intervention-based audit and Latent Vulnerability Score. Across three model families, constructed models showed **2.5–3.1× the vulnerability** of aligned bases at the targeted middle layer, exposing risks missed by behavioral evaluations.
- **CLEAR:** Co-led hidden-state gating of a safety LoRA adapter. Reduced HarmBench attack success from **32.3% to 0.5%** on Llama-3-8B-Instruct, with up to **7.1 percentage points higher GSM8K accuracy** than globally applied safety SFT or LoRA.
- **GANPO:** Designed latent adversarial regularization compatible with DPO and SimPO; improved robustness under distribution shift and high-entropy sampling with minor overhead. [Code]
- **MATCHA:** Built answer-conditioned adversarial probes and automated evaluation to expose inconsistent reasoning despite correct answers; studied transfer to black-box LLMs. [Code]

Robustness & Applied ML

RAMP/CURE: Multi-norm adversarial and certified training with gradient projection and bound alignment. **PyroAdapt:** Wildfire adaptation raised daily average precision from **21.62% to 24.57%** across California. **Federated domain adaptation:** Gradient projection and auto-weighted aggregation for healthcare and scarce target data.

Publications

* denotes equal contribution.

- Chengxiao Wang*, **Enyi Jiang***, Xiaojing Liao, Sanmi Koyejo. *CLEAR: Continuous Latent Adapter Routing for Utility-Preserving LLM Safety Alignment.* Trustworthy AI for Good Workshop @ NeurIPS 2026 (accepted). (Paper)
- **Enyi Jiang***, Anders Gjølbye*, Yibo Jacky Zhang, Sanmi Koyejo. *When Behavioral Safety Evaluation Fails: A Representation-Level Perspective.* Preprint, 2026; accepted by Stanford Causal Science Conference, Frontiers in AI Evaluation. (Paper) (Poster)
- **Enyi Jiang**, Wu Sun. *PyroAdapt: Adapting Wildfire Prediction under Spatial Heterogeneity and Temporal Shift.* Preprint, 2026. (Paper)
- **Enyi Jiang**, Yibo Jacky Zhang, Yinglun Xu, Andreas Haupt, Nancy Amato, Sanmi Koyejo. *Latent Adversarial Regularization for Offline Preference Optimization.* Preprint, 2026. (Paper)
- **Enyi Jiang**, Changming Xu, Nischay Singh, Tian Qiu, Gagandeep Singh. *Robust Answers, Fragile Logic: Probing the Decoupling Hypothesis in LLM Reasoning.* (TMLR 2026)
- **Enyi Jiang**, David S. Cheung, Gagandeep Singh. *Towards Generalized Certified Robustness with Multi-Norm Training.* (TMLR 2026)
- **Enyi Jiang**, Gagandeep Singh. *RAMP: Boosting Adversarial Robustness Against Multiple ℓ_p Perturbations for Universal Robustness.* (NeurIPS 2024)
- **Enyi Jiang***, Yibo Jacky Zhang*, Oluwasanmi O. Koyejo. *Principled Federated Domain Adaptation: Gradient Projection and Auto-Weighting.* (ICLR 2024)
- **Enyi Jiang**. *Federated Domain Adaptation for Healthcare.* University of Illinois Urbana-Champaign, 2023. (M.S. Thesis)
- **Enyi Jiang**, Haisong Zhang. *A New Coarse and Fine-grained Chinese Word Segmentation Benchmark.* Chinese patent application no. CN202210157392.7, 2022.
- Xuan Wang, Yingjun Guan, Weili Liu, Aabhas Chauhan, **Enyi Jiang**, Qi Li, David Liem, Dibakar Sigdel, John Caufield, Peipei Ping, Jiawei Han. *EVIDENCEMINER: Textual Evidence Discovery for Life Sciences.* ACL 2020, System Demonstrations. (Paper) (System demo)
- Shirui Luo, Jiahuan Cui, Madhu Vellakal, Jian Liu, **Enyi Jiang**, Seid Koric, Volodymyr Kindratenko. *Review and Examination of Input Feature Preparation Methods and Machine Learning Models for Turbulence Modeling.* 2019. (Paper)

Service

Reviewer: NeurIPS, ICLR, ICML, TMLR, CVPR, ECCV, AAAI, AISTATS, COLM, TPAMI.
Mentor, AI Alignment @ UIUC: Red-teaming/Adversarial and Interpretability streams.

Achievements

Second Place, UIUC Ashby Prize in Computational Science Hackathon	2022
Outstanding Graduate, Zhejiang University	2020
Dean's List, University of Illinois Urbana-Champaign	2019
Meritorious Winner, Mathematical Contest in Modeling	2018
Academic First-Class Scholarship, ZJU-UIUC Institute	2017, 2018
First-Class Student Award, Zhejiang University	2017